

# Part 1: Solving Problems with Search

# Lecture plan

- Python for AI
- Project 0: questions, comments?
- Agents that Plan
- Planning as Search
- Specific search algorithms
  - Some review, some new ones


# Why Python for AI?

- Object oriented, but without the annoyances of Java and C++
- Easy and fun to write code  
<https://www.learnpython.org/>
- Community and textbook support  
<https://wiki.python.org/moin/PythonForArtificialIntelligence>  
<http://aima.cs.berkeley.edu/python/readme.html>
- Advanced language features adapted from from LISP, Prolog:
  - functions as first-class citizens
  - List processing / lambda operators (from LISP)  
[http://www.u.arizona.edu/~erdmann/mse350/topics/list\\_comprehensions.html](http://www.u.arizona.edu/~erdmann/mse350/topics/list_comprehensions.html)

# Project 0: Python 101. Due Friday 1/20

- Goal: gentle introduction to Python language, common data structures
- Algorithm:
  1. Download files
  2. Follow tutorial (step-by-step)
  3. Read provided starting code
  4. Fill in function bodies
  5. Run autograder
  6. If all tests passed – submit. Else: goto 3.
- Questions, stuck? Piazza!
- **Python clinic:** Wednesday 1/18, Lab. (times TBD).

# Search




[Slides adapted from Dan Klein and Pieter Abbeel]


<http://ai.berkeley.edu>.]

# Today

- Agents that Plan Ahead
- Search Problems
- Uninformed Search Methods
  - Depth-First Search
  - Breadth-First Search
  - Iterative Deepening



# Agents that Plan

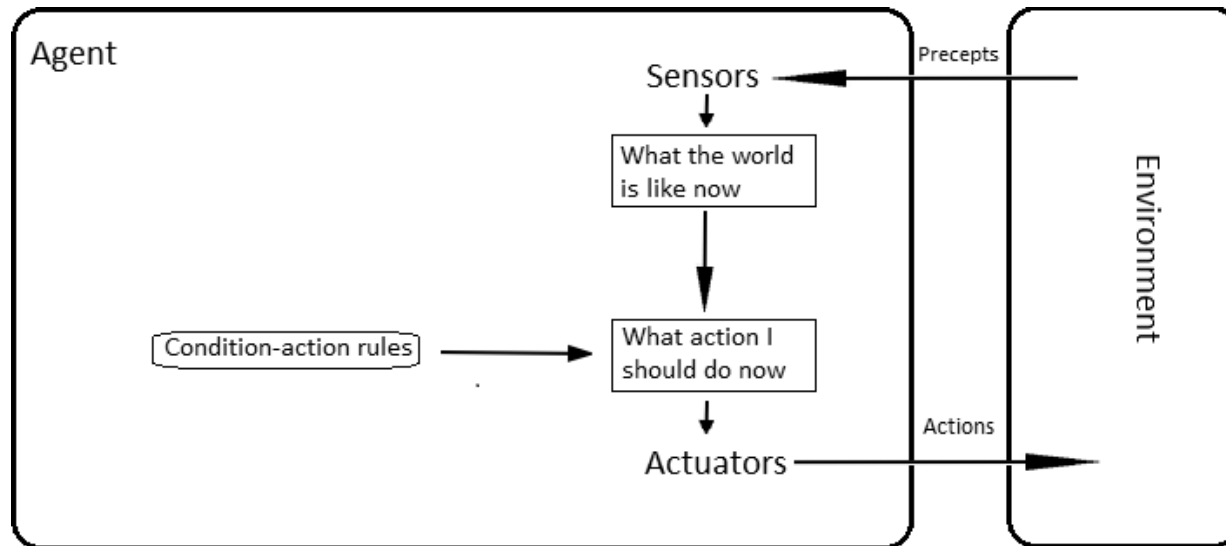


# Agent types

- Five basic types in order of increasing generality:
  - Table Driven agents
  - Simple reflex agents
  - Model-based reflex agents
  - Goal-based agents
    - Problem-solving agents
  - Utility-based agents
    - Can distinguish between different goals
  - Learning agents




# Reflex Agents



- Act only on the basis of the current percept.
- Based on the *condition-action rules*:  
if condition  
then action.
- Infinite loops are often unavoidable (can be fixed with some randomization)

# Reflex Agents


- Reflex agents:
  - Choose action based on current percept (and maybe memory)
  - May have memory or a model of the world's current state
  - Do not consider the future consequences of their actions
  - Consider how the world IS
- Can a reflex agent be rational?



Demo of reflex agent (optimal)  
Demo of reflex agent (stuck)

# Planning Agents

- Planning agents:
  - Ask “what if”
  - Decisions based on (hypothesized) consequences of actions
  - Must have a model of how the world evolves in response to actions
  - Must formulate a goal (test)
  - **Consider how the world **WOULD BE****
- Optimal vs. complete planning
- Planning vs. replanning



# Rational agents

- **Performance measure:** An objective criterion for success of an agent's behavior, e.g.,
  - Robot driver?
  - Chess-playing program?
  - Spam email classifier?
- **Rational Agent:** selects actions that is *expected* to **maximize its performance measure**,
  - given percept sequence
  - given agent's built-in knowledge
  - consider: how to maximize expected future performance, with only historical data?

# Task Environment

- Before we design an intelligent agent, we must specify its “task environment”:

## **PEAS:**

**P**erformance measure

**E**nvironment

**A**ctuators

**S**ensors

# PEAS: Example

- Example: Agent = robot driver in DARPA Challenge
  - Performance measure:
    - Time to complete course
  - Environment:
    - Road, obstacles
  - Actuators:
    - Steering wheel, accelerator, brake, signal, horn
  - Sensors:
    - Optical cameras, lasers, sonar, accelerometer, speedometer, GPS, odometer, engine sensors, ...



Stanley

# Goal-Based Agents: Search

- Five basic types in order of increasing generality:
  - Table Driven agents
  - Simple reflex agents
  - Model-based reflex agents
  - Goal-based agents
    - Problem-solving agents: solve problems by searching for solution
  - Utility-based agents
    - Can distinguish between different goals
  - Learning agents

# Search Problems







# Search Problems

- A **search problem** consists of:

- A state space




- A successor function  
(with actions, costs)



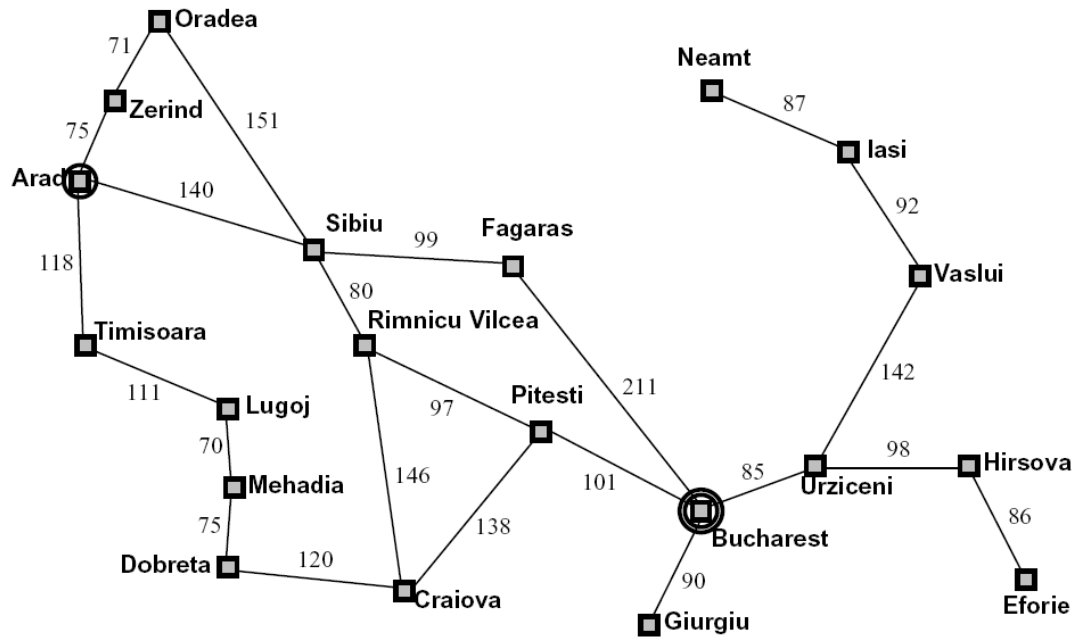
- A start state and a goal test

- A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state

# Search Problems Are Models




# Example: Traveling in Romania



- State space:
  - Cities
- Successor function:
  - Roads: Go to adjacent city with cost = distance
- Start state:
  - Arad
- Goal test:
  - Is state == Bucharest?
- Solution?

# What's in a State Space?

The **world state** includes every last detail of the environment




A **search state** keeps **only** the details needed for planning (abstraction).

Note: *Static info (e.g., walls) do not need to be stored for each state*


- Problem: Path-Finding
  - States:  $(x,y)$  location
  - Actions: NSEW
  - Successor: update location only
  - Goal test: is  $(x,y)=\text{END}$
- Problem: Eat-All-Dots
  - States:  $\{(x,y), \text{dot booleans}\}$
  - Actions: NSEW
  - Successor: update location and possibly a dot boolean
  - Goal test: dots all false

# State Space Sizes?

- World state:
  - Agent positions: 120
  - Food count: 30
  - Ghost positions: 12
  - Agent facing: NSEW
- How many
  - World states?  
 $120 \times (2^{30}) \times (12^2) \times 4$
  - States for path-finding?  
120
  - States for eat-all-dots?  
 $120 \times (2^{30})$




# Example: 8-queens problem



# State-Space problem formulation

- states? -any arrangement of  $n \leq 8$  queens  
-or arrangements of  $n \leq 8$  queens in leftmost  $n$  columns, 1 per column, such that no queen attacks any other.
- initial state? no queens on the board
- actions? -add queen to any empty square  
-or add queen to leftmost empty square such that it is not attacked by other queens.
- goal test? 8 queens on the board, none attacked.
- path cost? 1 per move



# Example: 8-puzzle

- states?
- initial state?
- actions?
- goal test?
- path cost?

|   |   |   |
|---|---|---|
| 7 | 2 | 4 |
| 5 |   | 6 |
| 8 | 3 | 1 |

Start State

|   |   |   |
|---|---|---|
|   | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State



# State Space Problem Formulation

- states? locations of tiles
- initial state? given
- actions? move blank left, right, up, down
- goal test? goal state (given)
- path cost? 1 per move


|   |   |   |
|---|---|---|
| 7 | 2 | 4 |
| 5 |   | 6 |
| 8 | 3 | 1 |

Start State

|   |   |   |
|---|---|---|
|   | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State


# Quiz: Safe Passage



- Problem: eat all dots while keeping the ghosts scared
- What does the state space have to specify?
  - (agent position, dot booleans, power pellet booleans, remaining scared time)


# State Space Graphs: 1

- **State space graph:** A mathematical representation of a search problem
  - Nodes are (abstracted) world configurations
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
  - Careful: if there are loops in this graph, keep track of visited states
- We can rarely build this full graph in memory (it's too big), but it's a useful idea




# State Space Graphs: 2

- State space graph: A mathematical representation of a search problem
  - Nodes are (abstracted) world configurations
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a search graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



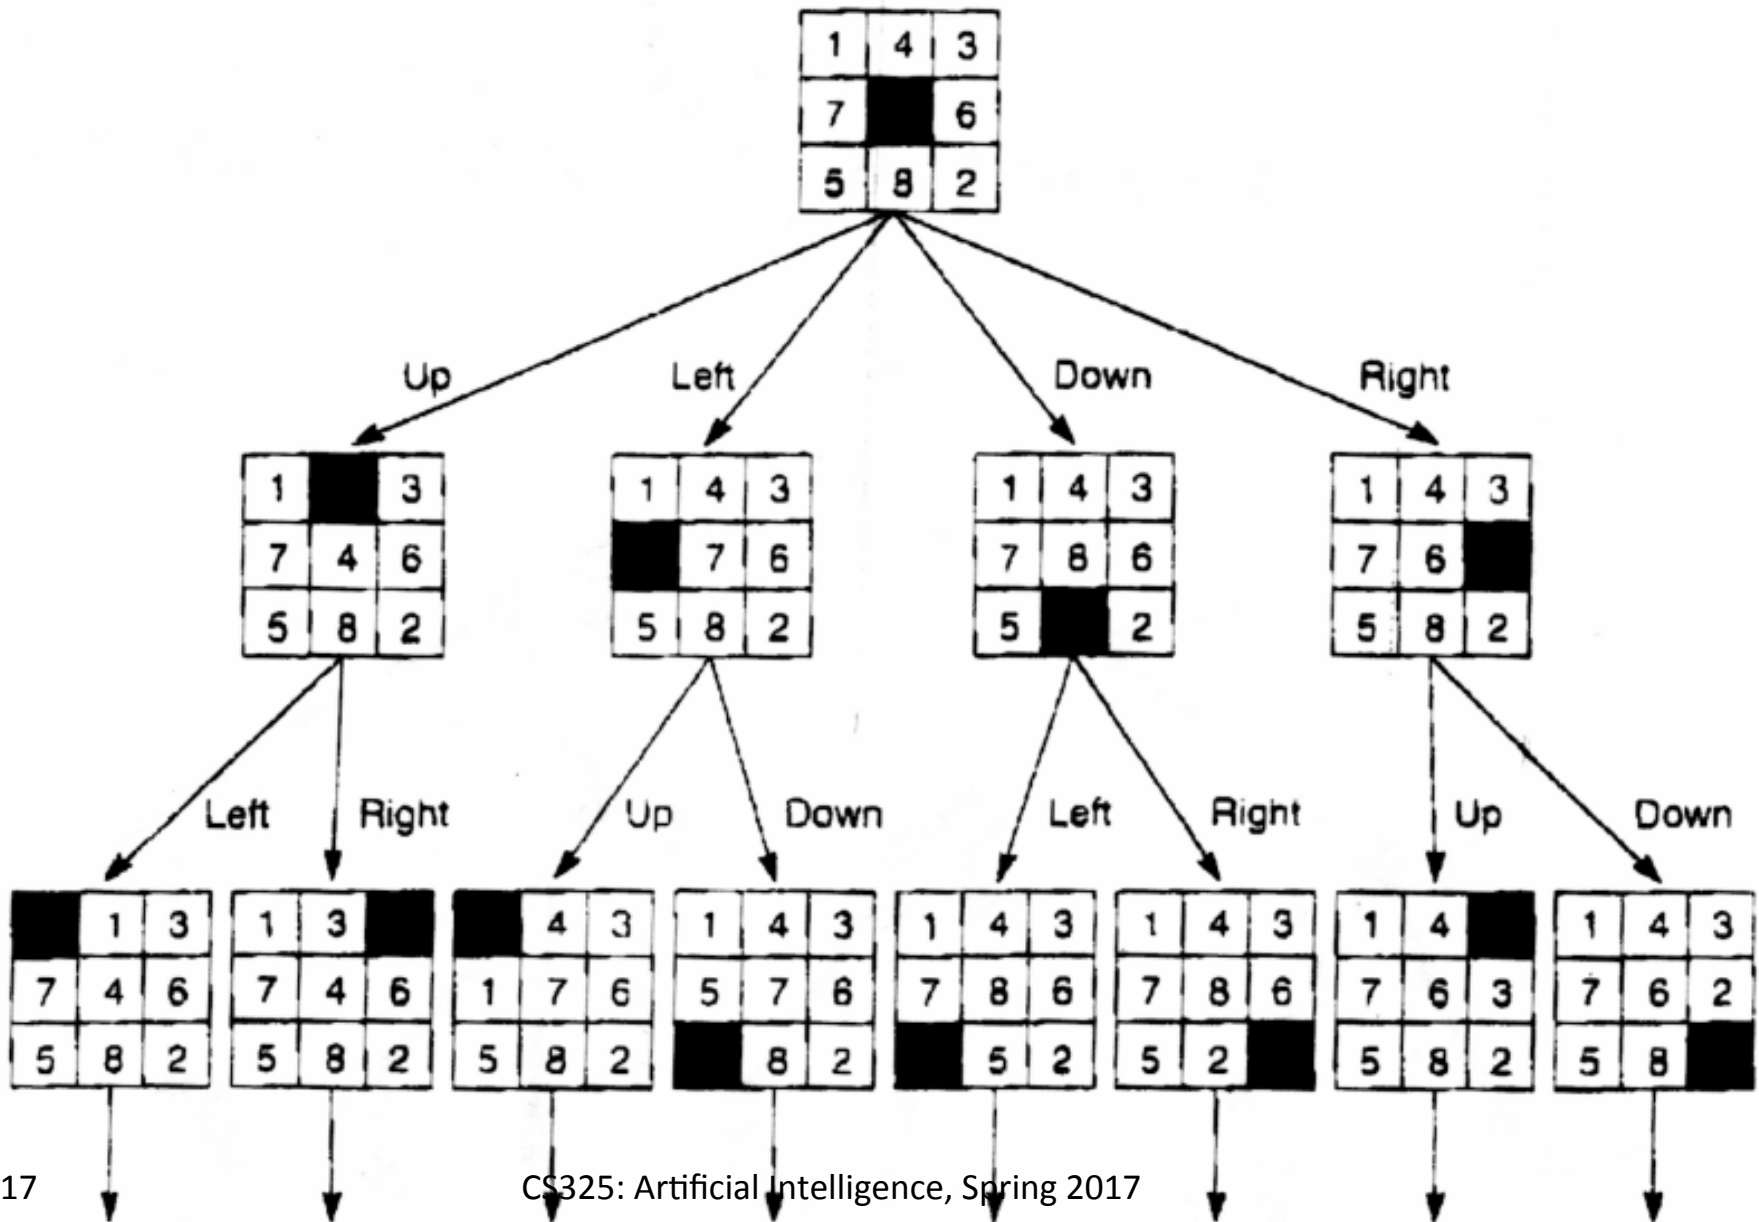
*Tiny search graph for a tiny search problem*

# Search Trees




- A search tree:
  - A “what if” tree of plans and their outcomes
  - The start state is the root node
  - Children correspond to successors
  - Nodes show states, but correspond to PLANS (paths from root to the state)
  - For most problems, we can never actually build the whole tree (too large)

# Search Tree for 8 puzzle problem




# State Space Graphs vs. Search Trees

State Space Graph




*Each NODE in the search tree is an entire PATH in the state space graph. We construct both on demand – and we construct as little as possible.*

Search Tree



# Quiz: State Space Graphs vs. Search Trees

Consider this 4-state graph:




How big is its search tree (from S)?



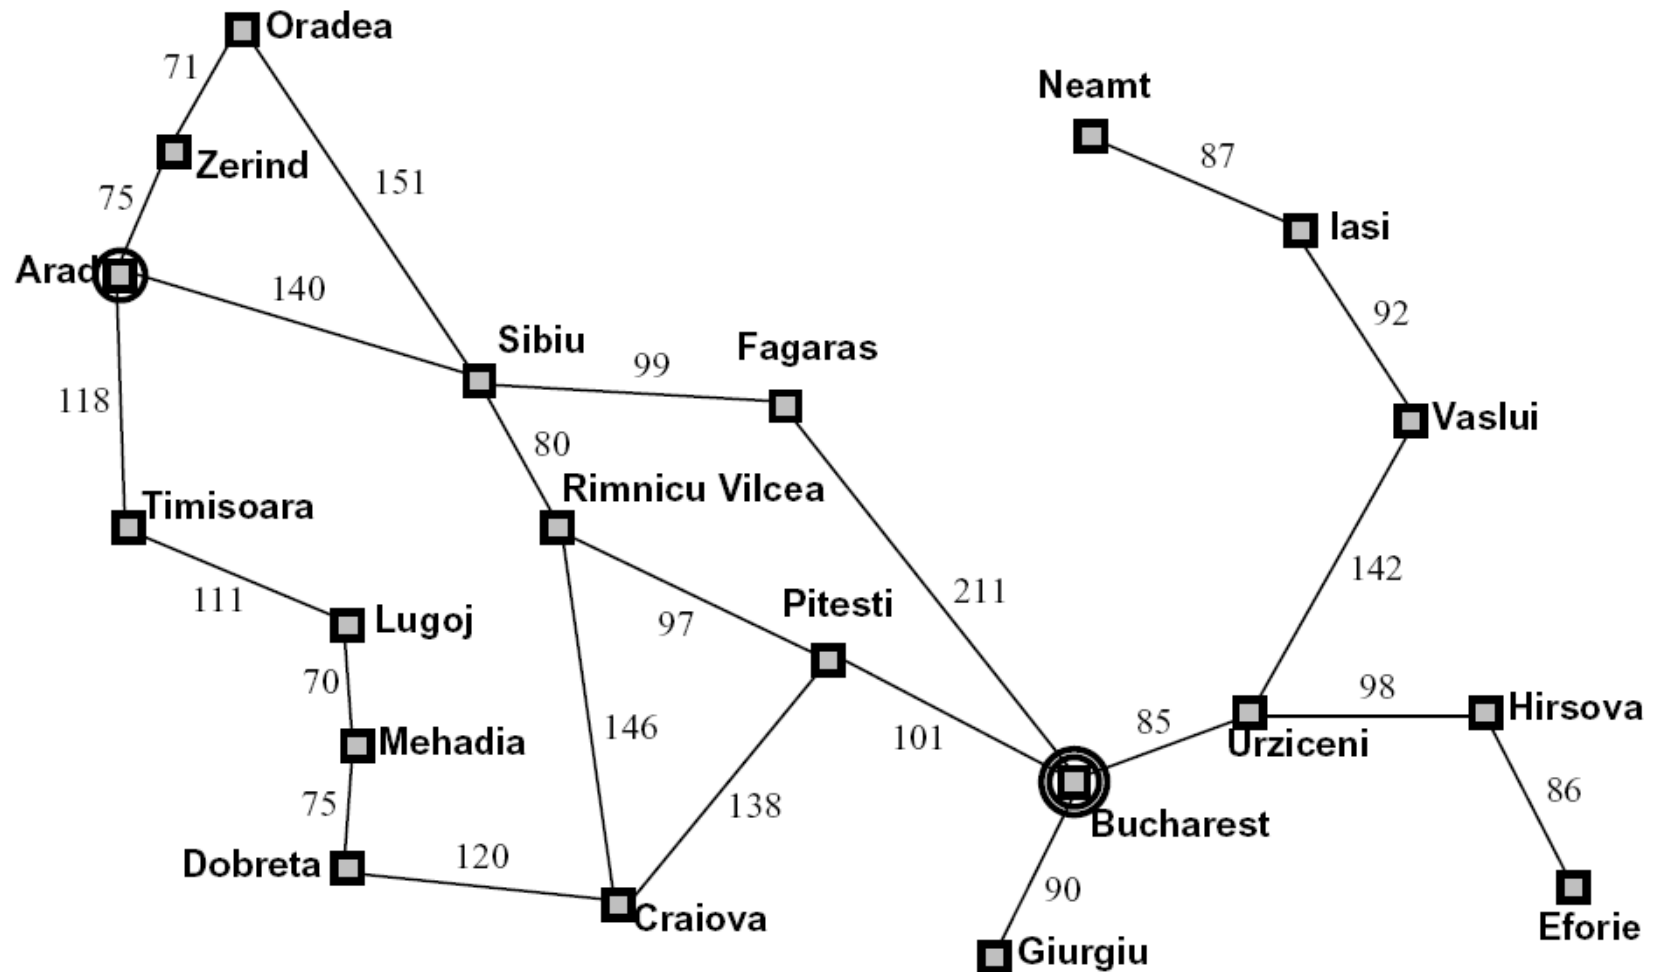
Problem: Lots of repeated structure in the search tree!




# Tree Search



# Search Example: Romania



# Searching with a Search Tree



- Search:
  - Expand out potential plans (tree nodes)
  - Maintain a **fringe** of partial plans under consideration
  - Try to expand as few tree nodes as possible




# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- Important ideas:
  - Fringe
  - Expansion
  - Exploration strategy
- Main question: which fringe nodes to explore?

# Example: Tree Search



[Example: on board]

# Why Search can be hard

Assuming  $b=10$ , 1000 nodes/sec, 100 bytes/node

| Depth of Solution | Nodes to Expand | Time          | Memory        |
|-------------------|-----------------|---------------|---------------|
| 0                 | 1               | 1 millisecond | 100 bytes     |
| 2                 | 111             | 0.1 seconds   | 11 kbytes     |
| 4                 | 11,111          | 11 seconds    | 1 megabyte    |
| 8                 | $10^8$          | 31 hours      | 11 giabytes   |
| 12                | $10^{12}$       | 35 years      | 111 terabytes |



$$P(40) \approx \frac{64!}{32! (8!)^2 (2!)^6} \approx 10^{43}.$$

# Sidebar: Search vs. Intuition



- Human chess grandmasters think “only” 3-5 moves ahead (Kasparov occasionally 12-14) but rely on patterns, intuition
- Deep blue and others: exhaustive search for optimal state/solution. Evaluates 100M positions/sec, vs. Kasparov 3 positions/sec

# Search Strategies

- A **search strategy** is defined by picking the order of node expansion (fringe exploration)
- Strategies are evaluated along the following dimensions:
  - **completeness**: does it always find a solution if one exists?
  - **time complexity**: number of nodes generated
  - **space complexity**: maximum number of nodes in memory
  - **optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
  - **b**: maximum branching factor of the search tree
  - **d**: depth of the least-cost solution
  - **m**: maximum depth of the state space (may be  $\infty$ )





# Depth-First Search




# Depth-First Search

*Strategy: expand  
a deepest node  
first*

*Implementation:  
Fringe is a LIFO  
stack*




# Search Algorithm Properties




# Search Algorithm Properties

- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?
- Cartoon of search tree:
  - $b$  is the branching factor
  - $m$  is the maximum depth
  - solutions at various depths
- Number of nodes in entire tree?
  - $1 + b + b^2 + \dots + b^m = O(b^m)$



# Depth-First Search (DFS) Properties

- What nodes DFS expand?
  - Some left prefix of the tree.
  - Could process the whole tree!
  - If  $m$  is finite, takes time  $O(b^m)$
- How much space does the fringe take?
  - Only has siblings on path to root, so  $O(bm)$
- Is it complete?
  - $m$  could be infinite, so only if we prevent cycles (more later)
- Is it optimal?
  - No, it finds the “leftmost” solution, regardless of depth or cost



# To Do:

- Enroll in Piazza, check for updates!
- Do the readings (Chapter 3 in R&N)
- Finish Project 0 (posted on website and Piazza)
- Next week: Project 1 will be assigned. hardness: 7
  - **Implication:** go through Python tutorials and do project 0 as soon as possible.