


# Solving Problems with Search: 4

# Today's Plan



- Project 1 comments and Q&A
- Review:  $A^*$
- Properties of  $A^*$  algorithm and heuristics

# Graph Search




# Tree Search: Extra Work!

- Failure to detect repeated states can cause exponentially more work.



# Graph Search

- In BFS, for example, we shouldn't bother expanding the circled nodes (why?)



# Graph Search: **Implementation**

- Idea: never **expand** a state twice
- How to implement:
  - Tree search + set of expanded states (“closed set”)
  - Expand the search tree node-by-node, but...
  - **Before expanding a node**, check to make sure its state has never been expanded before
  - If expanded: skip it, if new: add to closed set
- **Efficiency tip:** **store the closed set as a set**, not a list (**why?**)

# Graph Search Pseudo-Code

Why do you need graph search for Pacman?

```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      for child-node in EXPAND(STATE[node], problem) do
        fringe ← INSERT(child-node, fringe)
      end
    end
  end
```

Need this check to handle loops


# Project 1 Questions & Tips

- Use Piazza, read FAQ before posting questions:  
<https://piazza.com/emory/spring2017/cs325/>
- **Questions 1-4:** if you develop a correct solution for DFS, the rest will be very easy modifications
- Do not use shortcuts: use Node class or similar:  
<https://piazza.com/class/ixql4613j9k223>
- Questions 5-8: more fun/creative. Leave enough time.

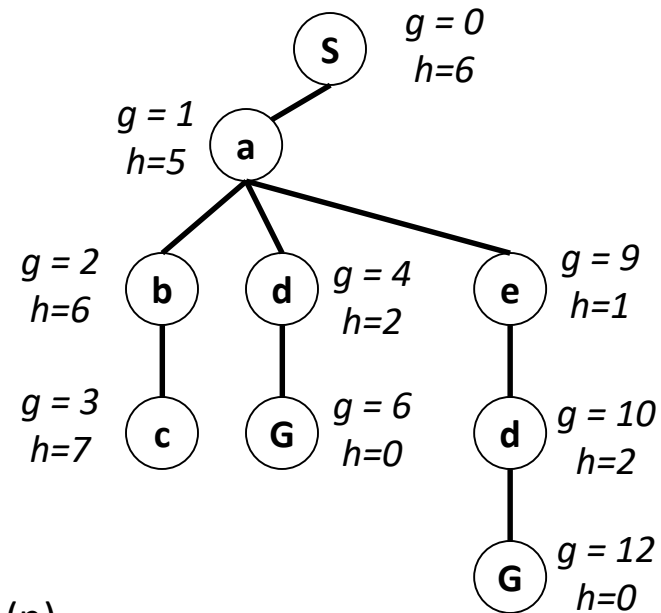


# A\* Review: $f(n) = \text{UCS} + \text{Heuristic}$

- Uniform-cost orders by path cost, or *backward cost*  $g(n)$
- Greedy orders by goal proximity, or *forward cost*  $h(n)$



- A\* Search orders by the sum:  $f(n) = g(n) + h(n)$



Example: Teg Grenager


# Admissible Heuristics

- A heuristic  $h$  is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where  $h^*(n)$  is the true cost to a nearest goal

- Examples:




- Coming up with admissible heuristics is most of what's involved in using A\* in practice.




# Properties of $A^*$

# Properties of A\*

Uniform-Cost





A\*



# UCS vs A\* Contours

- Uniform-cost expands equally in all “directions”
- A\* expands mainly toward the goal, but does hedge its bets to ensure optimality




[Demo: contours UCS / greedy / A\* empty (L3D1)]

[Demo: contours A\* pacman small maze (L3D5)]


# Video of Demo Contours (Empty) -- UCS

---




# Video of Demo Contours (Empty) -- Greedy

---




# Video of Demo Contours (Empty) – $A^*$

---







# Video of Demo Contours (Pacman Small Maze) – A\*




# Which Algorithm (1)?




# Which Algorithm (2)




# Which Algorithm (3)



# Which Algorithm (4)?



# Which Algorithm (5)



# Comparison: Summary



Greedy




Uniform Cost



A\*

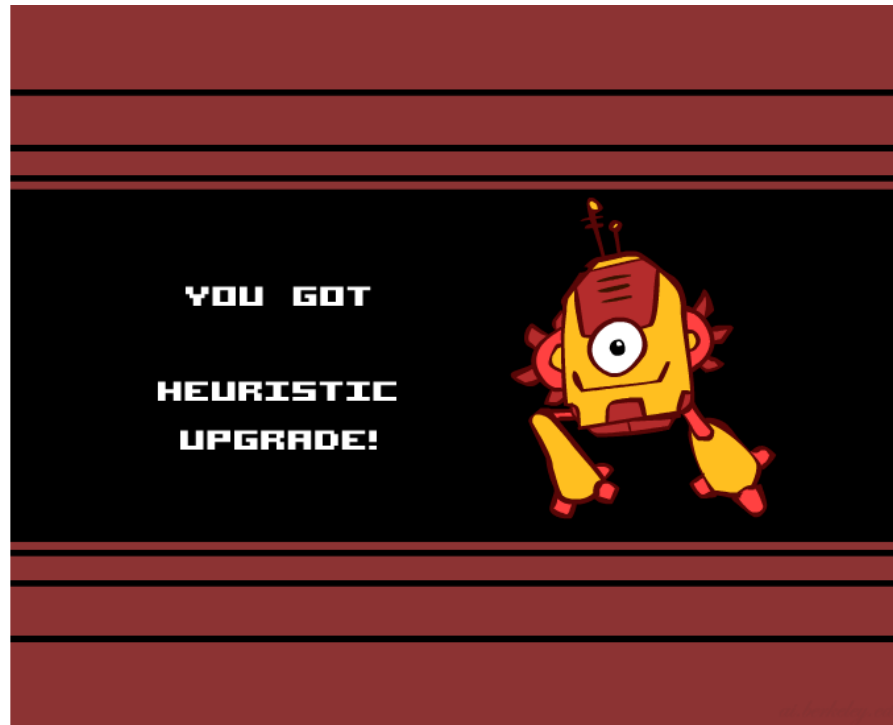
# A\* Applications

- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- ...






# Creating Heuristics



# Creating Admissible Heuristics

- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available




- Inadmissible heuristics are often useful too!

# Example: 8 Puzzle

7	2	4
5		6
8	3	1

Start State



Actions


	1	2
3	4	5
6	7	8

Goal State


- What are the states?
- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?

# 8 Puzzle I

- Heuristic:
- Is it admissible?
- $h(\text{start}) =$
- $h(\text{goal}) =$




Start State




Goal State

# 8 Puzzle: Tiles heuristic


- Heuristic: **Number of tiles misplaced**
- Why is it admissible?
- $h(\text{start}) = 8$
- This is a *relaxed-problem* heuristic



Start State



Goal State




Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	$3.6 \times 10^6$
TILES	13	39	227

Statistics from Andrew Moore


# 8 Puzzle II: **Manhattan** heuristic

- **Relaxation**: easier 8-puzzle where **any** tile could slide any direction at any time, ignoring other tiles?

- Total **Manhattan** distance from correct location



Start State



Goal State

- Is it admissible?


$$3 + 1 + 2 + \dots = 18$$

- $h(\text{start}) =$

Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
TILES	13	39	227
MANHATTAN	12	25	73

# 8 Puzzle III: **Oracle** heuristic

- How about using the *actual cost* as a heuristic?
  - Would it be admissible?
  - Would we save on nodes expanded?
  - What's wrong with it?



- With A\*: a **trade-off** between **quality of estimate** and **work per node**
  - As heuristics get closer to the true cost, you will **expand fewer nodes** but usually **do more work per node** to compute the heuristic itself

# Recap: Problem Relaxation

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then  $h_1(n)$  gives the shortest solution
- If the rules are relaxed so that a tile can move to **any adjacent square**, then  $h_2(n)$  gives the shortest solution



# Designing heuristics (cont'd)

E.g., for the 8-puzzle:

- $h_1(n)$  = number of misplaced tiles
- $h_2(n)$  = total Manhattan distance  
(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$
- $h_2(S) = ?$

# Heuristics: cont'd

E.g., for the 8-puzzle:

- $h_1(n)$  = number of misplaced tiles
- $h_2(n)$  = total Manhattan distance  
(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$  8
- $h_2(S) = ?$   $3+1+2+2+2+3+3+2 = 18$

Which is “better”  
–  $h_1$  or  $h_2$ ?

# Idea: Heuristic **dominance**

- If  $h_2(n) \geq h_1(n)$  for all  $n$  (both admissible, i.e.,  $<$  true cost)  
then  $h_2$  **dominates**  $h_1$

$h_2$  is better for search

- Typical search costs (average number of nodes expanded):
- $d=12$       IDS = 3,644,035 nodes  
     $A^*(h_1) = 227$  nodes  
     $A^*(h_2) = 73$  nodes
- $d=24$       IDS = too many nodes  
     $A^*(h_1) = 39,135$  nodes  
     $A^*(h_2) = 1,641$  nodes

# Example: Heuristics for Chess

- To select next move, must evaluate expected benefit of successor position:
  - **Value of the pieces** (count value of your pieces – value of opponents pieces)
  - **Space**: threatened/controlled space by you – space controlled by opponent
  - **Pawn** structure
  - ...
- Examples:
  - <https://www.quora.com/What-are-some-heuristics-for-quickly-evaluating-chess-positions>
  - <https://chessprogramming.wikispaces.com/Killer+Heuristic>

# Example: Heuristics for Motion Planning

- Robot motion: many moving (body) parts
- What's the most efficient way to accomplish goal?

<https://www.youtube.com/watch?v=dSwDZmvtGZY>

# Example: Machine Translation

- 1. Translate words from source to target
- 2. Choose the “more likely” translation among candidates
- $h(t)$  = count of phrase seen in target language
- What could go wrong...

# Example: Machine Translation

- $h(t)$  = count of phrase seen in target language




# Designing Heuristics

- A good heuristic is:
  - ✓ Admissible (optimistic)
  - Consistent (non-decreasing)
  - ✓ “Accurate”




# Trivial Heuristics, Dominance

- Dominance:  $h_a \geq h_c$  if
$$\forall n : h_a(n) \geq h_c(n)$$
- Heuristics form a **semi-lattice**:
  - Max of admissible heuristics is admissible
$$h(n) = \max(h_a(n), h_b(n))$$
- Trivial heuristics
  - Bottom of lattice is the zero heuristic (what does this give us?)
  - Top of lattice is the exact heuristic




# A\* Graph Search Gone Wrong?


State space graph



Search tree




# Consistency of Heuristics




- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal  
 $h(A) \leq \text{actual cost from A to G}$
  - Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc  
 $h(A) - h(C) \leq \text{cost(A to C)}$
- Consequences of consistency:
  - The f value along a path never decreases  
 $h(A) \leq \text{cost(A to C)} + h(C)$
  - A\* graph search is optimal

# Optimality of A\* Graph Search



# Optimality of A\* Graph Search


- Sketch: consider what A\* does with a **consistent** heuristic:
  - **Fact 1:** In tree search, A\* expands nodes in increasing total  $f$  value ( $f$ -contours)
  - **Fact 2:** For every state  $s$ , nodes that reach  $s$  optimally are expanded before nodes that reach  $s$  suboptimally
  - Result: A\* graph search is optimal



??? Uhm... we already proved that for  
**admissible heuristics???**

# Optimality (2): Tree vs. Graph Search

- **Tree search:**
  - A\* is optimal if heuristic is **admissible**
  - UCS is a special case ( $h = 0$ )
- **Graph search:**
  - A\* optimal if heuristic is **consistent**
  - UCS optimal ( $h = 0$  is consistent)
- Consistency implies admissibility
- In general, **most natural admissible heuristics tend to be consistent**, especially if from relaxed problems



# A\*: Summary



# A\*: Summary

- A\* uses both backward costs and (estimates of) forward costs
- A\* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems

