

Local Search

Today's Plan

- Project 1 comments and Q&A
- Iterative Improvement for CSPs → Local Search

Project 1 Tips

- Use Piazza, read FAQ before posting questions:
<https://piazza.com/class/ixql4613j9k223?cid=8>
- **Questions 1-4:** if you develop a correct solution for DFS, the rest will be easy modifications
- **Run autograder** after *every* question. Until you pass all the test cases, assume your code has bugs.
- It's OK to (re) submit multiple times, only latest attempt is graded.


Tips for Project 1 (cont'd)

- Problems 5-8 depend on code in 1-4. Get that right (and tested) first, before moving on!
- P5/Corners problem: must visit all corners in *single* path
 - Implications for search tree, state info to update
- Heuristics for p6-8: start simple. For extra credit, think back to graph traversal algorithms from cs323.

Reminder: CSPs


- CSPs:

- Variables
- Domains
- Constraints
 - Implicit (provide code to compute)
 - Explicit (provide a list of the legal tuples)
 - Unary / Binary / N-ary



- Goals:

- Here: find any solution
- Also: find all, find best, etc.



Backtracking Search

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

Why more efficient to use a recursive alg (not iterative + Fringe like regular search)?

Improving Backtracking

- General-purpose ideas give huge gains in speed
 - ... but it's all still NP-hard
- **Filtering**: Can we detect inevitable failure early?
- **Ordering**:
 - Which variable should be assigned next? (MRV)
 - In what order should its values be tried? (LCV)
- Many other ideas (e.g., problem structure)




Iterative Algorithm for CSPs: MinConflicts


- Local search methods typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
 - Take an assignment with unsatisfied constraints
 - Operators *reassign* variable values
 - No fringe!** Live on the edge.
- Greedy algorithm: While not solved,
 - Variable selection: randomly select **any** conflicted variable
 - Value selection: min-conflicts heuristic:
 - Choose a value that violates the fewest constraints
 - I.e., **hill climb** with $h(n) = \text{total number of violated constraints}$



Example: 4-Queens



$n = 5$





[Board]

- States: 4 queens in 4 columns ($4^4 = 256$ states)
- Operators: move queen in column
- Goal test: no attacks
- Evaluation: $c(n)$ = number of attacks

Performance of Min-Conflicts


- Given random initial state, can solve n-queens in almost constant time for arbitrary n with high probability (e.g., $n = 10,000,000$)!
- The same appears to be true for any randomly-generated CSP *except* in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$



Summary: CSPs

- CSPs are a special kind of search problem:
 - States are partial assignments
 - Goal test defined by
- Basic solution: backtracking search
- Speed-ups:
 - Ordering
 - Filtering
- Iterative min-conflicts is often effective in practice




Google's Deep Mind Wins at Go

- <http://deepmind.com/alpha-go.html>
- “Simple” game, but difficult to master
- Orders of magnitude more states than chess (x “googol”)
- Traditional search (A^*) not feasible
- Google solution:
 - Local search w/ restarts (this lecture)
 - “deep learning” to estimate state values (instead of heuristics) ← observing human players
 - <https://googleblog.blogspot.com/2016/01/alphago-machine-learning-game-go.html>




Searching Large (or Infinite) Spaces

- Too many states to explore using A* or variants
 - Large or infinite branching factor (continuous)
- “Reasonable” solution is good enough
- Local search idea: **start with initial guess and incrementally improve**



Local Search

- Tree search keeps unexplored alternatives on the fringe (ensures completeness)
- Local search: improve a single option until you can't make it better (no fringe!)
- New successor function: local changes




- Generally much faster and more memory efficient (but incomplete and suboptimal)

Hill Climbing

evaluation

$h(s)$: Estimate of distance from a peak (smaller is better)



OR: Height Defined by Evaluation Function (greater is better)



Hill Climbing: Trade-offs

- What's bad about this approach?
 - Complete?
 - Optimal?
- What's good about it?



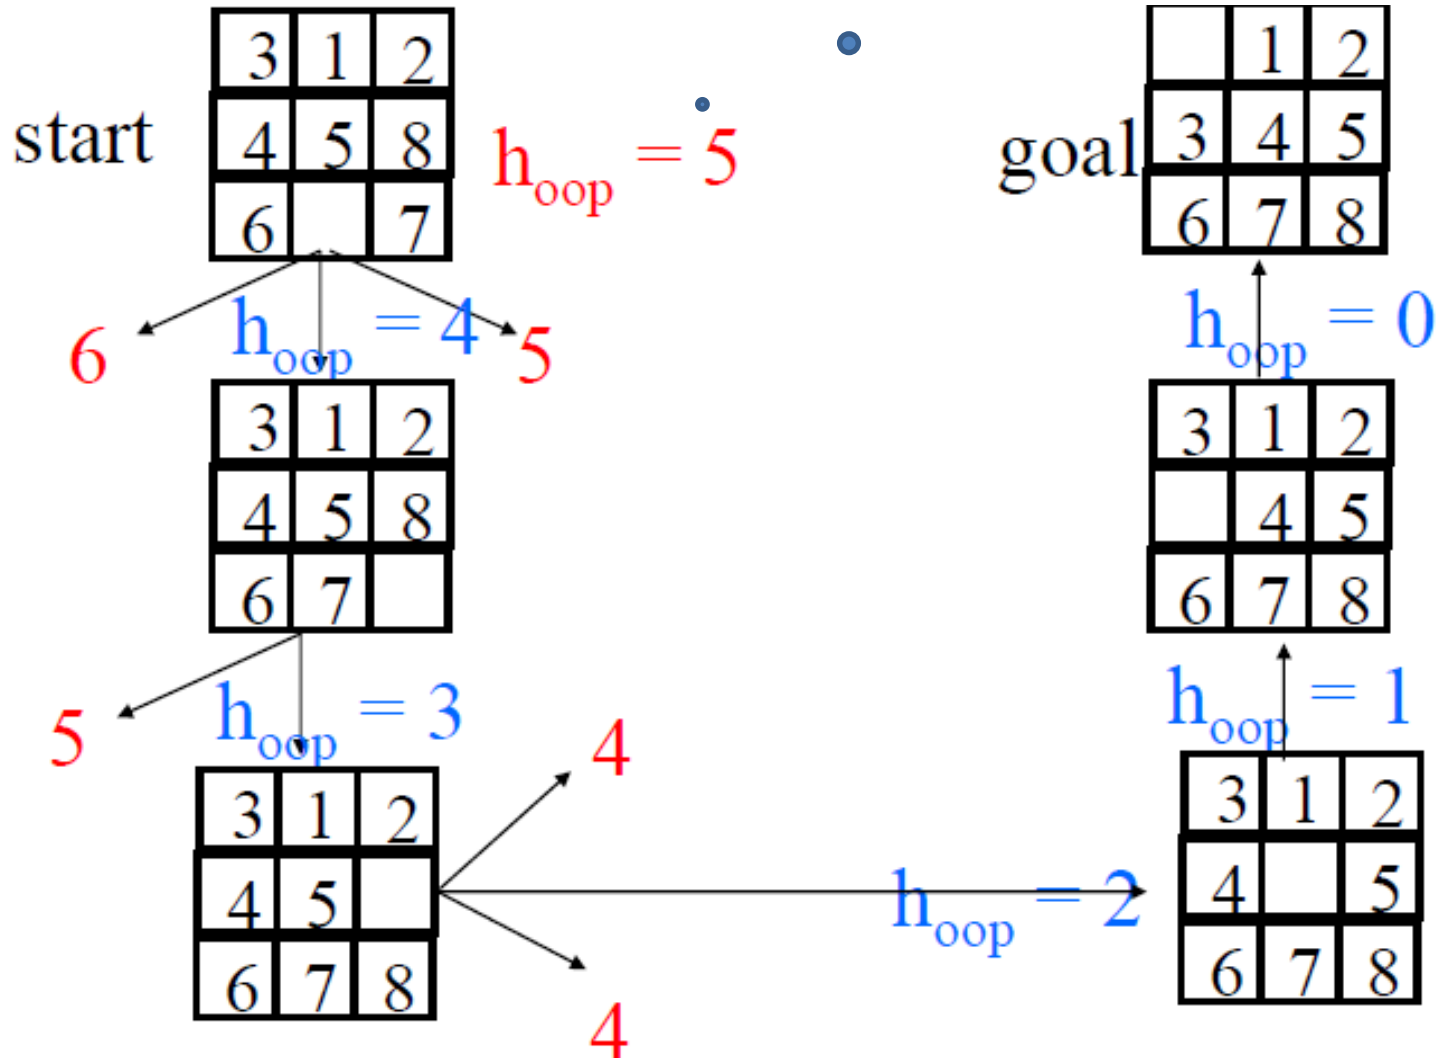
Hill Climbing: Algorithm

- I. **While** (\exists uphill points):
 - Move in the direction of increasing evaluation function f
- II. **Let** $s_{next} = \arg \max_s f(s)$, s a successor state to the current state n
 - If $f(n) < f(s)$ then move to s
 - Otherwise halt at n
- **Properties:**
 - Terminates when a peak is reached.
 - Does not look ahead of the immediate neighbors of the current state.
 - Chooses randomly among the set of best successors, if there is more than one.
 - Doesn't *backtrack*, since it doesn't remember where it's been
- **a.k.a. *greedy local search***


"Like climbing Everest in thick fog with amnesia"

Toy Example: Tiles


h = out of place (oop) tiles



Hill Climbing Diagram



Hill Climbing Quiz




Starting from X, where do you end up ?

Starting from Y, where do you end up ?

Starting from Z, where do you end up ?

Drawbacks of Hill Climbing

- **Local Maxima:** peaks that aren't the highest point in the space
- **Plateaus:** the space has a broad flat region that gives the search algorithm no direction (random walk)




- **Ridges:** dropoffs to the sides; steps to the North, East, South and West may go down, but a step to the NW may go up.

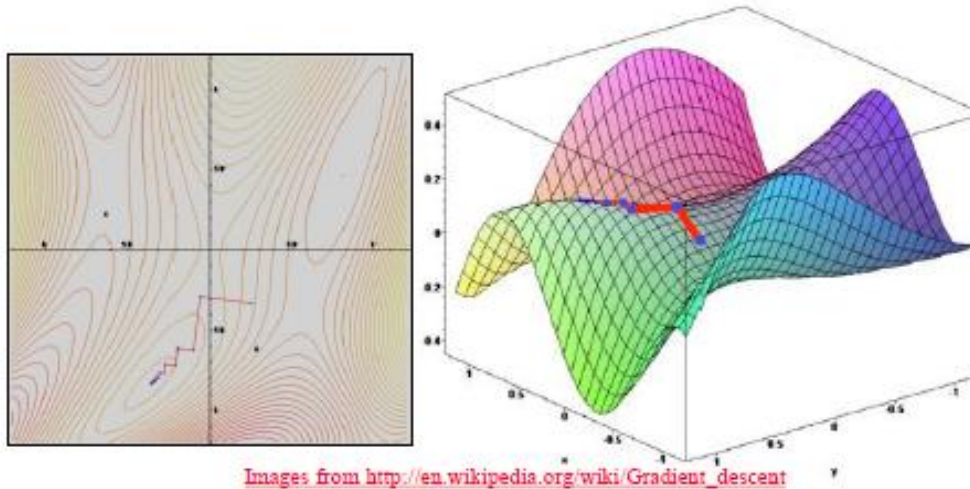


“Easy” Problems: Convex Surface

- No local maxima (only 1 peak)
- Hill climbing works great
- Can we make it faster?



Gradient Descent (Steepest Descent)




- Gradient descent procedure for finding the $\arg_x \min f(x)$
 - choose initial x_0 randomly
 - repeat
 - $x_{i+1} \leftarrow x_i - \eta f'(x_i)$
 - until the sequence $x_0, x_1, \dots, x_i, x_{i+1}$ converges
- Step size η (eta) is small (perhaps 0.1 or 0.05)

<http://vis.supstat.com/2013/03/gradient-descent-algorithm-with-r/>

Gradient Ascent vs. Newton-Ralphston

- A reminder of Newton's method from Calculus:
$$x_{i+1} \leftarrow x_i - \eta f'(x_i) / f''(x_i)$$
- Newton's method uses 2nd order information (the second derivative, or, curvature) to take a more direct route to the minimum.
- The second-order information is more expensive to compute, but converges quicker.




Contour lines of a function
Gradient descent (green)
Newton's method (red)

[Image from http://en.wikipedia.org/wiki/Newton's_method_in_optimization](http://en.wikipedia.org/wiki/Newton's_method_in_optimization)

(this and previous slide from Eric Eaton)




Problem: Non-Convex Surfaces

- Realistic problems:
 - Many local suboptimal maxima
 - Easy to get trapped



Problem: Non-Convex Surfaces

- Realistic problems:
 - Many local suboptimal maxima
 - Easy to get trapped
- Examples:




Solving the Problems

- Allow backtracking (What happens to complexity?)
- Stochastic hill climbing: choose at random from uphill moves, using steepness for a probability
- Random restarts: “If at first you don’t succeed, try, try again.”
- Several moves in each of several directions, then test
- Jump to a different part of the search space

Random Restart (Monte-Carlo methods)

- Idea: restart hill climbing algorithm from random start configurations
- Repeat N times.
- If reasonable sampling of space, w high prob will find global max

- In the end: Some problem spaces are great for hill climbing and others are terrible.



Monte Carlo Descent


- 1) $S \leftarrow$ initial state
- 2) Repeat k times:
 - a) If GOAL?(S) then return S
 - b) $S' \leftarrow$ successor of S picked at random
 - c) if $h(S') \leq h(S)$ then $S \leftarrow S'$
 - d) else
 - $Dh = h(S') - h(S)$
 - with probability $\sim \exp(-Dh/T)$, where T is called the “temperature”,
do: $S \leftarrow S'$ [Metropolis criterion]
- 3) Return failure

Simulated Annealing (2)

- Variant of hill climbing (maximize value)
- Tries to **explore** enough of the search space **early on**, so that the optimal solution is less sensitive to the start state
- May make some **downhill moves** before finding a good way to move uphill.

Simulated Annealing (2)

- Comes from the physical process of annealing in which substances are raised to high energy levels (melted) and then cooled to solid state.



- The probability of moving to a higher energy state, instead of lower is

$$p = e^{(-\Delta E/kT)}$$

where ΔE is the positive change in energy level, T is the temperature, and k is Boltzmann's constant.

Simulated Annealing: Intuition

- At the beginning, the temperature is high.
- As the temperature becomes lower
 - kT becomes lower
 - $\Delta E/kT$ gets bigger
 - $(-\Delta E/kT)$ gets smaller
 - $e^{(-\Delta E/kT)}$ gets smaller
- As the process continues, the probability of a downhill move gets smaller and smaller.
- ΔE is the change in the value of the objective function.
- Need an **annealing schedule**, which is a sequence of values of T : T_0, T_1, T_2, \dots

Simulated Annealing Algorithm


- Idea: Escape local maxima by allowing downhill moves
 - But make them rarer as time goes on

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to “temperature”
  local variables: current, a node
                   next, a node
                   T, a “temperature” controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\textit{next}] - \text{VALUE}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E/T}$ 
```

Implementation: “Select successor with probability p ”

- Select *next* with probability p



- Generate a random number
- If it's $\leq p$, select *next*

Simulated Annealing Properties

- Theoretical guarantee:
 - Stationary distribution: $p(x) \propto e^{-\frac{E(x)}{kT}}$
 - If T decreased slowly enough, will converge to optimal state!
- Is this an interesting guarantee?
- Sounds like magic, but reality is reality:
 - The more downhill steps you need to escape a local optimum, the less likely you are to ever make them all in a row
 - People think hard about *ridge operators* which let you jump around the space in better ways

Simulated Annealing Properties

- At a fixed “temperature” T , state occupation probability reaches the Boltzman distribution: https://en.wikipedia.org/wiki/Boltzmann_distribution

$$p_i = \frac{e^{-\varepsilon_i/kT}}{\sum_{j=1}^M e^{-\varepsilon_j/kT}}$$


- If T is decreased **slowly enough** (very slowly), the procedure will reach the best state.
- **Slowly enough** has proven too slow for some researchers who have developed alternate schedules.

Simulated Annealing Schedules

- Acceptance criterion and cooling schedule

if ($\text{delta} \geq 0$) accept

else if ($\text{random} < e^{\text{delta} / \text{Temp}}$) accept, else reject /* $0 \leq \text{random} \leq 1$ */



Initially temperature is very high (most bad moves accepted)

Temp slowly goes to 0, with multiple moves attempted at each temperature

Final runs with temp=0 (always reject bad moves) greedily “quench” the system

Simulated Annealing Applications

- Basic Problems


- Traveling salesman
- Graph partitioning
- Matching problems
- Graph coloring
- Scheduling

<http://toddwschneider.com/posts/traveling-salesman-with-simulated-annealing-r-and-shiny/>



- Engineering

- VLSI design
 - Placement
 - Routing
 - Array logic minimization
 - Layout
- Facilities layout
- Image processing
- Code design in information theory
- Chemistry: molecular structure:



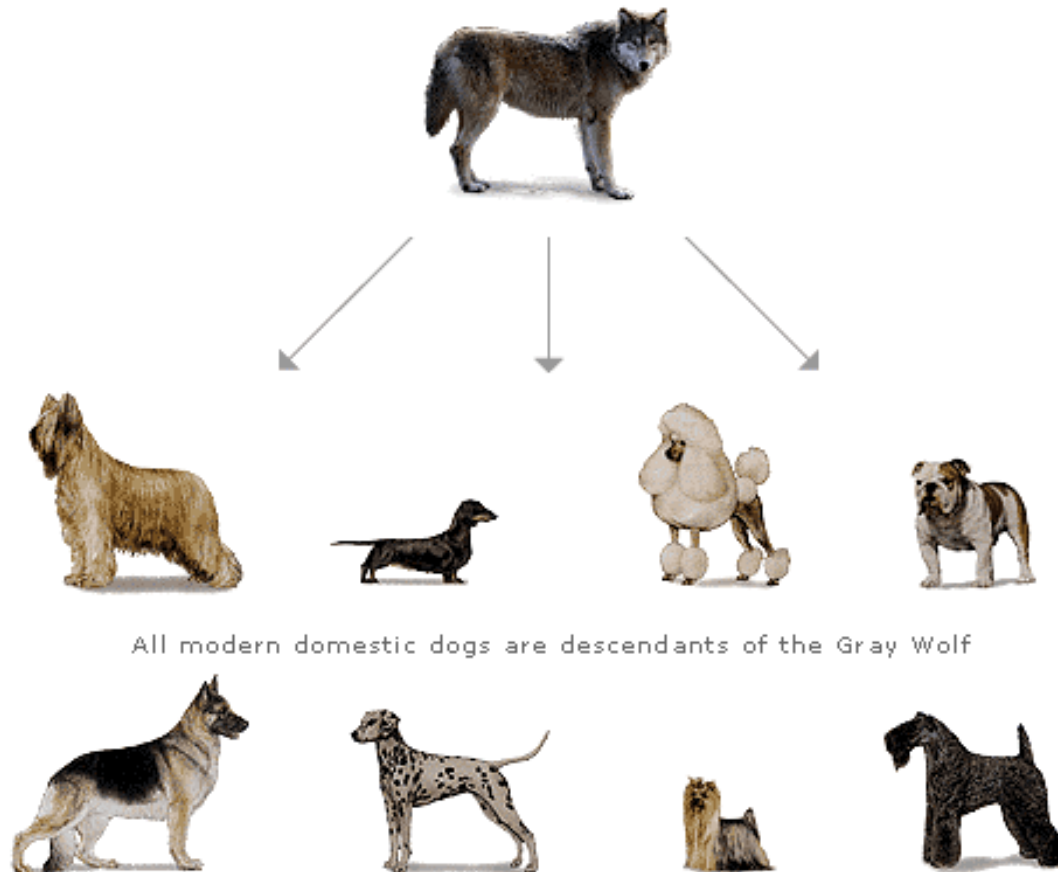
<http://www.sciencedirect.com/science/article/pii/S0166128097001954>

Local Beam Search

- Keeps more previous states in memory
 - Simulated annealing just kept one previous state in memory.
 - This search keeps k states in memory.
 - randomly generate k initial states
 - if any state is a goal, terminate
 - else, generate all successors and select best k
 - repeat


Explore Multiple Moves in Parallel?

- Evolution-inspired computation



Genetic Algorithms

- Start with k random states (the *initial population*)
- 2. New states are generated by either
 - 1. “*Mutation*” of a single state or
 - 2. “*Sexual Reproduction*”: (combining) two *parent states* (selected proportionally to their *fitness*) – *crossover*
- Encoding used for the “*genome*” of an individual strongly affects the behavior of the search
- Similar (in some ways) to stochastic beam search




Genetic Algorithm

- Given: population **P** and fitness-function **f**
- repeat
 - **newP** \leftarrow empty set
 - for $i = 1$ to **size(P)**
 - $x \leftarrow \text{RandomSelection}(P, f)$
 - $y \leftarrow \text{RandomSelection}(P, f)$
 - $child \leftarrow \text{Reproduce}(x, y)$
 - if (small random probability) then $child \leftarrow \text{Mutate}(child)$
 - add $child$ to **newP**
 - **P** \leftarrow **newP**
- until some individual is fit enough or enough time has elapsed
- return the best individual in **P** according to **f**

Example: Drive Train Design

- **Genes for:**
 - Number of Cylinders
 - RPM: 1st- \rightarrow 2nd
 - RPM 2nd- \rightarrow 3rd
 - RPM 3rd- \rightarrow Drive
 - Rear end gear ratio
 - Size of wheels
- **A chromosome specifies a full drive train design**



GA: Reproduction

- Reproduction by **crossover** selects genes from two parent chromosomes and creates two new offspring.
- To do this, randomly choose a crossover point (perhaps none).
- For child 1, everything before this point comes from the first parent and everything after from the second parent.
- Crossover looks like this (| is the crossover point):

Chromosome 1 11001 | 00100110110

Chromosome 2 10011 | 11000011110

Offspring 1 11001 | 11000011110

Offspring 2 10011 | 00100110110



GA: Mutation

- Mutation randomly changes genes in the new offspring.
- For binary encoding we can switch randomly chosen bits from 1 to 0 or from 0 to 1.

Original offspring 1101111000011110

Mutated offspring 1100111000001110

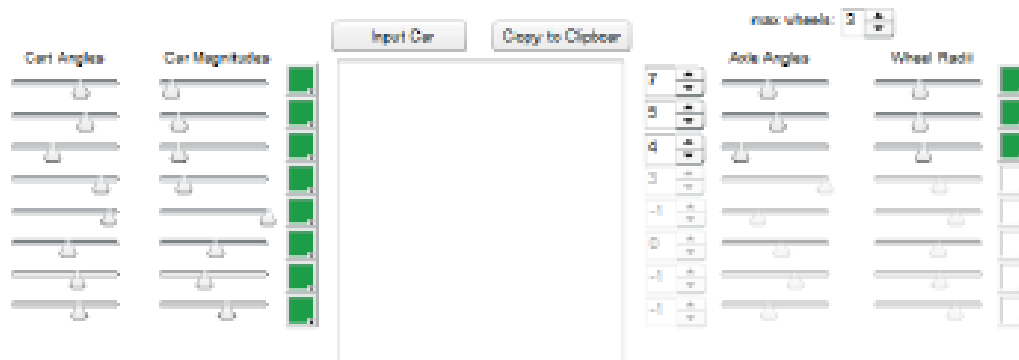


GA Example: Cont'd

BoxCar 2D

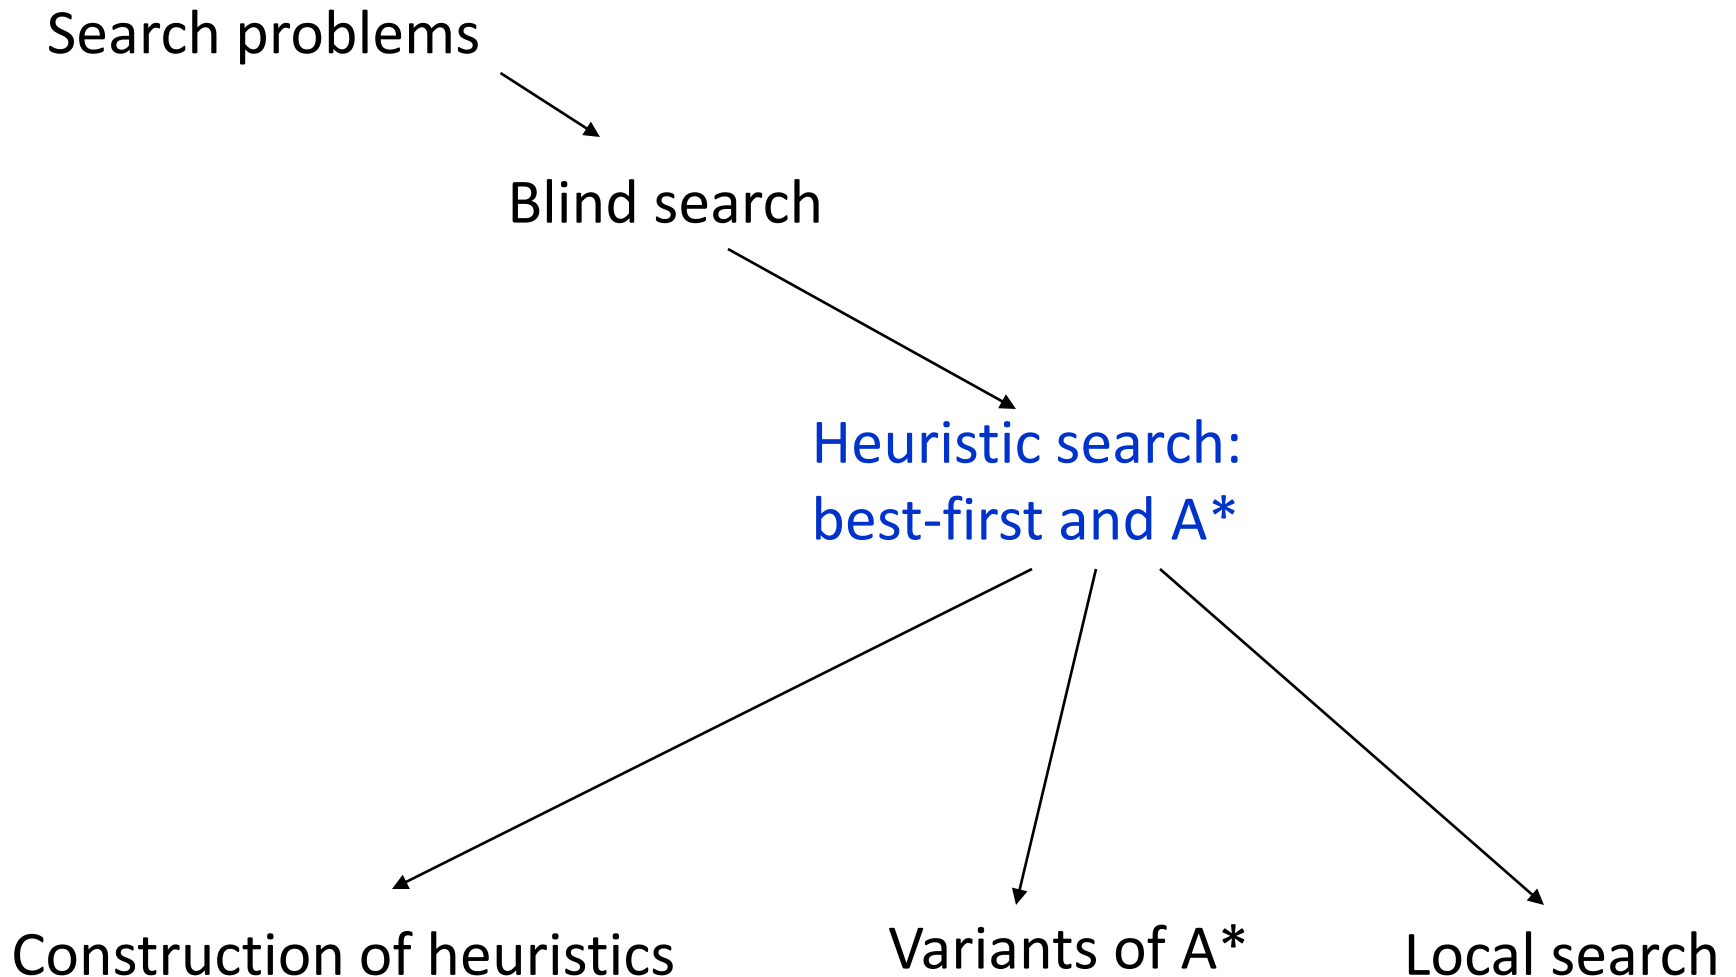
[Home](#) | [Designer](#) | [News](#) | [FAQ](#) | [The Algorithm](#) | [Versions](#) | [Contact](#)

Derp Bike Designer



<http://boxcar2d.com/>

Summary: Search Techniques So Far




Quick Review/Quiz

- What is a difference between a game state and search node?
- What do we lose if our A^* heuristic not admissible?
- Why can't we use A^* for “large” problems?

When to Use Search Techniques?

- 1) The search space is small, and
 - No other technique is available, or
 - Developing a more efficient technique is not worth the effort
- 2) The search space is large, and
 - No other available technique is available, and
 - There exist “good” heuristics

Thursday: Adversarial Search!



SCORE: 0

Announcements

- Reminder: **project 1 due Tomorrow (Wed) at 8pm**
 - Submit on Blackboard (enabled last week)
- Normal office hours tomorrow (Wed) 3-4, and Thursday 4-5.
- Thursday: start adversarial search;
Project 2 will be assigned, due in ~1.5 weeks
- Next week: Instructor away next Tuesday, Rafi Haque will continue adversarial search.