


Adversarial Search: More on Eval Functions, ExpectiMax

With slides from
Dan Klein, Percy Liang, Luke Zettlemoyer


Multiplayer Games: 1 vs. All

- MAX vs. multiple MIN players




- How does the algorithm change?


Multiplayer Game Tree: 1 Max, 2 MINs




Multiplayer Game Tree: depth 1




Multiplayer Game Tree: depth 2




Multiplayer MiniMax: depth 2




Multiplayer MiniMax: depth 2




Multiplayer MiniMax: depth 2



Multiplayer MiniMax: depth 2



Multiplayer MiniMax: depth 2



MiniMax Algorithm: Multi-Min version

```
function MINIMAX-DECISION(state) returns an action  
  inputs: state, current state in game  
   $v \leftarrow \text{MAX-VALUE}(\text{state})$   
  return the action in SUCCESSORS(state) with value  $v$ 
```

```
function MAX-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for  $a, s$  in SUCCESSORS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, 1))$   
  return  $v$ 
```

Note:
MAX=Agent1
MIN1=Agent2
MIN2=Agent3
....
MINK=AgentN

```
function MIN-VALUE(state, agentIndex) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for  $a, s$  in SUCCESSORS(agentIndex, state) do  
    if (agentIndex  $\geq$  numAgents) then  
       $v_{temp} = \text{MAX-VALUE}(s)$   
    else #another ghost plays  
       $v_{temp} = \text{MIN-VALUE}(s, \text{agentIndex}+1)$   
     $v \leftarrow \text{MIN}(v, v_{temp})$   
  return  $v$ 
```

problem: collusion possible

- Previous slide (standard minimax analysis) assumes that each player operates to maximize only their own utility
- In practice, players **could make alliances**
 - Ex: C strong, A and B both weak
 - May be best for A and B to attack C rather than each other
- If game is not zero-sum (i.e., $\text{utility}(A) = - \text{utility}(B)$) then alliances can be useful even with 2 players
 - e.g., both cooperate to maximum the sum of the utilities
- Ignore this, **assume non-cooperative games**


Minimax Algorithm: Analysis

- Complete depth-first exploration of the game tree
- Performance:
 - Max depth = d , b legal moves at each point
 - E.g., Chess: $d \sim 100$, $b \sim 35$

Criterion	Minimax
Time	$O(\mathbf{b}^m)$ ☹️
Space	$O(bm)$ 😊

Resource Limits

- Problem: In realistic games, cannot search to leaves!
- Solution: ? Hint: 2 complementary approaches?
 1. ?
 2. ?
- Example:
 - Suppose we have 100 seconds, can explore 10K nodes / sec
 - So can check 1M nodes per move
 - α - β reaches about depth 8 – decent chess program
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Use iterative deepening for an anytime algorithm



Alpha-Beta Pseudocode

inputs: *state*, current game state

α , value of best alternative for MAX on path to *state*

β , value of best alternative for MIN on path to *state*

returns: *a utility value*

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ )
  if TERMINAL-TEST(state) then
    return UTILITY(state)
   $v \leftarrow -\infty$ 
  for  $a, s$  in SUCCESSORS(state) do
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$ 
    if  $v \geq \beta$  then return  $v$ 
     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ )
  if TERMINAL-TEST(state) then
    return UTILITY(state)
   $v \leftarrow +\infty$ 
  for  $a, s$  in SUCCESSORS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```

Alpha-Beta Pruning Properties

- This pruning has **no effect** on final result at the root
- Values of intermediate nodes might be wrong!
 - but, they are bounds
- Good child ordering improves effectiveness of pruning
- With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!
 - Full search of, e.g. chess, is still hopeless...

AlphaBeta Experiment


- `time python pacman.py -p MinimaxAgent -l minimaxClassic -a depth=4 -q`
- `time python pacman.py -p AlphaBetaAgent -l minimaxClassic -a depth=4 -q`

More Alpha-Beta Pruning Examples

- Generic game tree visualization:
<http://homepage.ufp.pt/jtorres/ensino/ia/alfabeta.html>
- Example Play: Connect Four:
<https://gimu.org/connect-four-js/>

(Better) Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g. $f_1(s) = (\text{num white queens} - \text{num black queens})$, etc.

Construction of an Evaluation Function

- Usually a weighted sum of “features”:

$$e(s) = \sum_{i=1}^n w_i f_i(s)$$

- Features may include
 - Number of pieces of each type
 - Number of possible moves
 - Number of squares controlled

More Examples of Evaluation Functions

- Reversi

- Number squares held?

- Better: number of squares held that **cannot** be flipped

- Prefer valuable squares

- NxN array $w[i,j]$ of position values

- Highest value: corners, edges

- Lowest value: next to corner or edge

- $s[i,j] = +1$ player, 0 empty, -1 opponent

$$score = \sum_{i,j} w[i,j]s[i,j]$$

Eval Function Approximation



Key idea: parameterized evaluation functions

$\text{Eval}(s; \mathbf{w})$ depends on weights $\mathbf{w} \in \mathbb{R}^d$

Feature vector: $\phi(s) \in \mathbb{R}^d$

$$\phi_1(s) = K - K'$$


$$\phi_2(s) = Q - Q'$$

...

Linear evaluation function:

$$\text{Eval}(s; \mathbf{w}) = \mathbf{w} \cdot \phi(s)$$

Chess Example: Revisited



Example: chess

$\text{Eval}(s) = \text{material} + \text{mobility} + \text{king-safety} + \text{center-control}$

$\text{material} = 10^{100}(K - K') + 9(Q - Q') + 5(R - R') +$
 $3(B - B' + N - N') + 1(P - P')$

$\text{mobility} = 0.1(\text{num-legal-moves} - \text{num-legal-moves}')$

...


More Evaluation Functions

- Chess:
 - eval(s) =
 - w1 * material(s) +
 - w2 * mobility(s) +
 - w3 * king safety(s) +
 - w4 * center control(s) + ...
 - In practice MiniMax improves accuracy of heuristic eval function
 - But one can construct pathological games where more search hurts performance!
(Nau 1981)

Idea: End-Game Databases


- Ken Thompson - all 5 piece end-games
- Lewis Stiller - all 6 piece end-games
 - Refuted common chess wisdom: many positions thought to be ties were really forced wins -- 90% for white
 - Is perfect chess a win for white?

The MONSTER



White wins in 255 moves
(Stiller, 1991)

Evaluation for Pacman



What features would be good for Pacman?

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$


Pacman Eval Functions

- Use a linear combination of features.

$$\text{Eval}(\text{state}) = f1 * w1 + f2 * w2$$

- Example:
 - f1: distance to nearest food
w1: ?
 - f2: distance to nearest ghost
w2: ?

Pacman Eval Functions Demo



- Alpha-beta pruning, simple eval function:

```
python pacman.py -p AlphaBetaAgent -a depth=4 -l minimaxClassic -q
```

```
python pacman.py -p AlphaBetaAgent -a depth=4 -l smallClassic -q
```

- Alpha-beta, better eval function:

```
python pacman.py -p AlphaBetaAgent -a depth=4,evalFn=better -l minimaxClassic -q
```

```
python pacman.py -p AlphaBetaAgent -a depth=4,evalFn=better -l smallClassic -q
```


Expected Minimax (ExpectiMax)

$$v = \sum_{\text{chance nodes}} P(n) \times \text{Minimax}(n)$$

$$3 = 0.5 \times 4 + 0.5 \times 2$$

Interleave chance nodes
with min/max nodes

Again, the tree is constructed
bottom-up





Idea: Maximum Expected Utility

- Why should we average utilities? Why not minimax?
- Principle of maximum expected utility: an agent should choose the action which **maximizes its expected utility, given its knowledge**
 - General principle for decision making
 - Often taken as the definition of rationality
 - We'll see this idea over and over in this course!
- Let's decompress this definition...


Reminder: Probabilities (More Later)

- A **random variable** represents an event whose outcome is unknown
- A **probability distribution** is an assignment of weights to outcomes
- Example: Traffic on freeway
 - Random variable: T = whether there's traffic
 - Outcomes: T in {none, light, heavy}
 - Distribution: $P(T=\text{none}) = 0.25$, $P(T=\text{light}) = 0.50$, $P(T=\text{heavy}) = 0.25$
- Some laws of probability (more later):
 - Probabilities are always non-negative
 - Probabilities over all possible outcomes sum to one
- As we get more evidence, probabilities may change:
 - $P(T=\text{heavy}) = 0.25$, $P(T=\text{heavy} \mid \text{Hour}=8\text{am}) = 0.60$
 - We'll talk about methods for reasoning and updating probabilities later




Reminder: Expectations

- The expected value of a function of a random variable is the average, weighted by the probability distribution over outcomes
- Example: How long to get to the airport?



Time: 20 min 30 min 60 min
 x + x + x
Probability: 0.25 0.50 0.25



35
min



ExpectiMax: Revised Pseudocode

```
def value(s)
```

```
    if s is a max node return maxValue(s)
```

```
    if s is an exp node return expValue(s)
```

```
    if s is a terminal node return evaluation(s)
```

```
def maxValue(s)
```

```
    values = [value(s') for s' in successors(s)]
```


```
    return max(values)
```

```
def expValue(s)
```

```
    values = [value(s') for s' in successors(s)]
```


```
    weights = [probability(s, s') for s' in successors(s)]
```

```
    return expectation(values, weights)
```





What Probabilities to Use?

- In expectimax search, we have a probabilistic model of how the opponent (or environment) will behave in any state
 - Model could be a simple uniform distribution (roll a die)
 - Model could be sophisticated and require a great deal of computation
 - We have a chance node for any outcome out of our control: opponent or environment
 - The model might say that adversarial actions are likely!
- For now, assume each chance node magically comes along with probabilities that specify the distribution over its outcomes



Having a probabilistic belief about another agent's action does not mean that the agent is flipping any coins!

Expectimax Pruning?



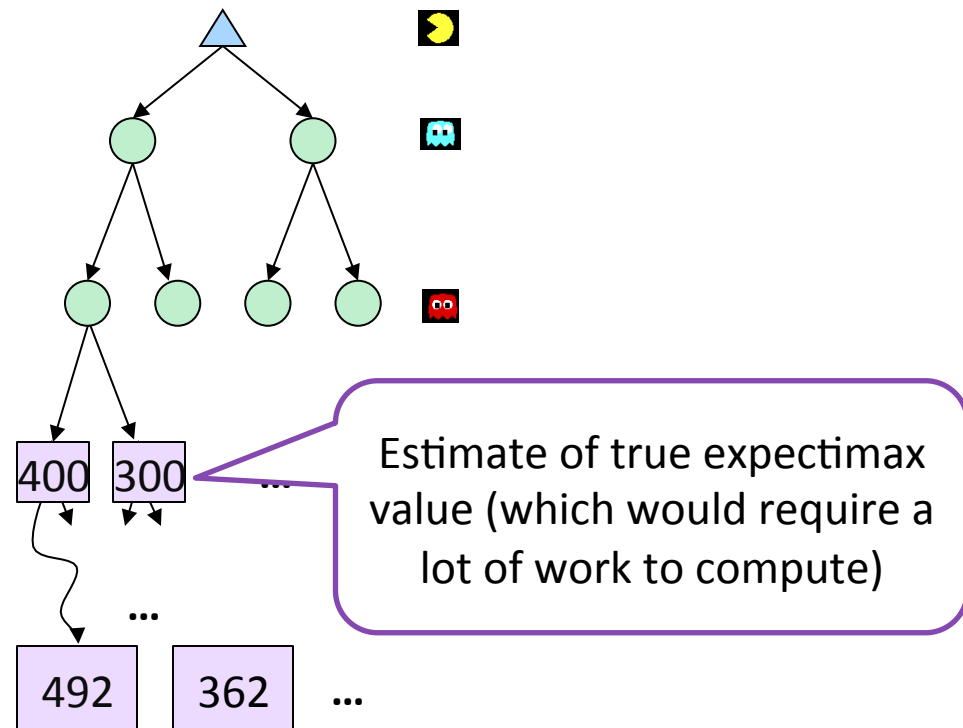
Theoretically possible, but:

- Either accept that with some probability we will discard optimal solution
- OR: no guarantee that pruning will save much time


First algorithm (I know of) was published in 2009, makes a lot of assumptions:

<http://ieeexplore.ieee.org/document/5286476/>

Depth-Limited Expectimax



What Utilities to Use?



- For worst-case minimax reasoning, terminal function scale doesn't matter
 - We just want better states to have higher evaluations (get the ordering right)
 - We call this **insensitivity to monotonic transformations**
- For average-case expectimax reasoning, we need *magnitudes* to be meaningful


Expectimax for Pacman

- Notice that we stepped away from thinking that the ghosts are trying to minimize pacman's score
- Instead, they are now a part of the environment
- Pacman has a belief (distribution) over how they will act
- **Quiz:** Can we model MiniMax as special case of ExpectiMax?

ExpectiMax for Pacman

Results from playing 5 games


	Minimizing Ghost	Random Ghost
Minimax Pacman	Won 5/5 Avg. Score: 493	Won 5/5 Avg. Score: 483
Expectimax Pacman	Won 1/5 Avg. Score: -303	Won 5/5 Avg. Score: 503



Pacman does depth 4 search with an eval function that avoids trouble
Minimizing ghost does depth 2 search with an eval function that seeks Pacman

Quiz: Informed Probabilities

- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?



- **Answer: Expectimax!**
 - To figure out EACH chance node's probabilities, you have to run a simulation of your opponent
 - This kind of thing gets very slow very quickly
 - Even worse if you have to simulate your opponent simulating you...
 - ... except for minimax, which has the nice property that it all collapses into one game tree

Project 2: Assigned, Due Wed 2/15

- <http://www.mathcs.emory.edu/~eugene/cs325/p2/>
 - Use Piazza for discussions
 - Questions?